

Mining Constrained Regions of Interest: An optimization approach

Alexandre Dubray Guillaume Derval Siegfried Nijssen Pierre Schaus

Table of contents

1. Introduction
2. PopularRegion
3. Our method
4. Experiments

Introduction


Motivations

- The amount of spatiotemporal data is exploding (smartphone applications, sports devices, fleet management, etc.)
- There is a need to process more efficiently these data
- Rewrite the raw trajectories (GPS points) as sequence of *Regions of Interest* (ROI)
- Multiple applications:
 - Trajectory pattern mining
 - Next location prediction
 - Urban management
 - ...

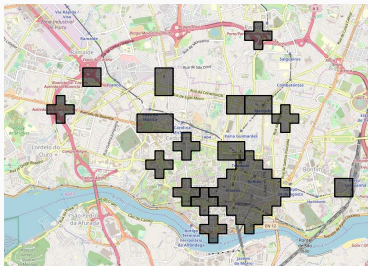
The general approach

1. Divide the map with a $N \times M$ grid.
2. Assign a density value to each cell. A cell is dense if its density is above a threshold. Multiple choice for the density:
 - Number of trajectory that passes in the cell
 - Number of trajectory that stays at least X minutes in the cell
 - etc.
3. Express the ROI as an aggregation of dense cells

Example of ROIs




(a) Initial set of dense cells




(b) Solution found by our method

PopularRegion


Execution of the algorithm




Execution of the algorithm




Execution of the algorithm




Execution of the algorithm




Execution of the algorithm




Execution of the algorithm




Execution of the algorithm



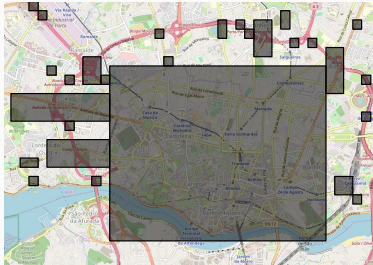
Execution of the algorithm



Result of the algorithm



(a) Initial set of dense cells



(b) Solution with 5% min average density

Advantages and disadvantages

- Scalable
- Intuitive and good results for most configurations


But...

- No formalization of the output
- Only rectangular regions
- Does not easily accept background knowledge
- Easy to create pathological input

Our method

ROIs as an encoder

- The ROIs encode the dense status of the cells
- Example of encoding with two rectangles
 - 1 dense cells is not covered
 - 4 non-dense cells are covered
 - The encoding makes 5 errors
- We prefer encoding with fewer errors



Formalization of the problem (1)

Some notations:

- Let \mathcal{G} be the grid, \mathcal{G}^* the set of dense cells, \mathcal{S} a set of candidates and θ the minimum density threshold
- d_i (resp. u_i) is the number of dense (resp. non-dense) cells covered by the candidate $R_i \in \mathcal{S}$
- K is the number of ROIs we want to find

A first optimization model

$$\text{minimize } |\mathcal{G}^*| + \overbrace{\sum_{R_i \in \mathcal{S}} x_i \cdot \underbrace{(u_i - d_i)}_{\text{errors of the candidate}}}^{\text{errors of the model}}$$

subject to

$$\begin{aligned} \sum_{R_i \in \mathcal{S}} x_i &\leq K \\ \sum_{R_i \in \mathcal{S} | c \in R_i} x_i &\leq 1 \quad \forall c \in \mathcal{G} \\ x_i &\in \{0, 1\} \quad \forall R_i \in \mathcal{S} \end{aligned}$$

Formalization of the problem (2)

In practice how to set the K ? Use the Minimum Description Length Principle!

- Let $Sol \subseteq \mathcal{S}$ be a valid selection of candidates
- Length of the errors (2 integers per cell):

$$L(\mathcal{G} \mid Sol) = 2(|\mathcal{G}^*| + \sum_{R_i \in Sol} (u_i - d_i))$$

- Length of the model:

$$L(Sol) = \sum_{R_i \in Sol} size(R_i)$$

- Minimum Description Length principle says that the best solution is:


$$\arg \min_{Sol \in \mathcal{S}} L(\mathcal{G} \mid Sol) + L(Sol) = 2|\mathcal{G}^*| + \sum_{R_i \in Sol} (2(u_i - d_i) + size(R_i))$$

The final optimization model


$$\begin{aligned} & \text{minimize} \quad \sum_{R_i \in \mathcal{S}} x_i \cdot \overbrace{(2(u_i - d_i) + \text{size}(R_i))}^{\text{Contribution to the description length}} \\ & \text{subject to} \\ & \quad \sum_{R_i \in \mathcal{S} \mid c \in R_i} x_i \leq 1 \quad \forall c \in \mathcal{G} \\ & \quad x_i \in \{0, 1\} \quad \forall R_i \in \mathcal{S} \end{aligned}$$

Example

- $L(\mathcal{S}) = 4 + 4 = 8$
- $L(\mathcal{G} \mid \mathcal{S}) = 2 \cdot (4 + 1) = 10$
- Total length of this model is $8 + 10 = 18$




- $L(\mathcal{S}) = 4 + 4 = 8$
- $L(\mathcal{G} \mid \mathcal{S}) = 2 \cdot (2 + 2) = 8$
- Total length of this model is $8 + 8 = 16$




Example with circles

- $L(\mathcal{S}) = 4 + 4 = 8$
- $L(\mathcal{G} \mid \mathcal{S}) = 2 \cdot (4 + 1) = 10$
- Total length of this model is $8 + 10 = 18$



- $L(\mathcal{S}) = 4 + 3 = 7$
- $L(\mathcal{G} \mid \mathcal{S}) = 2 \cdot (2 + 3) = 10$
- Total length of this model is $7 + 10 = 17$



1. Generate the set of candidates \mathcal{S} (e.g. enumerate all distinct rectangle on the grid)
 - Candidate can have any shape
 - Compute their contribution to the description length
 - Apply *intra-ROI* constraints to filter the candidate set
2. Solve the optimization model
 - Model *inter-ROI* constraints with linear constraints in the ILP
 - Solve the ILP, the binary decision variables give the set of ROIs

Experiments

- Two versions of our method
 - With only rectangular regions
 - With rectangular and circular regions
- Showing results on Kaggle taxis dataset (≈ 1.6 million trajectories)
- Comparing with PopularRegion¹ and OPTICS² (when clustering the dense cells)

¹Fosca Giannotti et al. "Trajectory pattern mining". In: *SIGKDD*. 2007.

²Mihael Ankerst et al. "OPTICS: ordering points to identify the clustering structure". In: *ACM Sigmod record* (1999).

Execution time

Minimum density threshold	2%			5%		
	100	150	200	100	150	200
Grid side size						
Number of dense cells ($ \mathcal{G}^* $)	571	597	537	230	178	137
Number of ILP candidates	23 814	7 779	3 399	2 880	1 232	434
ILP optimization time (s)	4.328	0.464	0.109	0.113	0.044	0.029
<i>PopularRegion</i> run time (s)	0.003	0.005	0.006	0.002	0.003	0.004
OPTICS run time (s)	0.209	0.222	0.200	0.084	0.065	0.051

Description Length

- For high density threshold, number of errors becomes similar
- ILP-based methods produce smaller models
- Overall the Description Length is inferior for ILP-based methods




Figure 3: Encoding of the errors




Figure 4: Encoding for the models

Robustness to noise

- Start from a 100×100 grid
- Move every element of the trajectories to a new cell with a probability p
- Choose the new cell randomly in a square of size 10 around the initial cell
- Compute the solution from the noisy grid and compute the *precision*, *recall* and *F1-measure*.




Figure 5: Recall and precision




Figure 6: F1-measure

Conclusion and Future work

What we did:

- We propose an optimization model to find K ROIs from trajectory data
- Our method is more flexible than specific method since it accepts a wide range of constraints
- The runtime of the ILP becomes reasonable as long as there is not too much candidates
- Everything is Open Source, see <https://github.com/AlexandreDubray/mining-roi>

The next steps:

- Get rid of the grid
- Use the density information (instead of just dense/not dense)
- Provide support for more complex constraints